
AURsec

A blockchain aproach to securing software packages

Bennett Piater & Lukas Krismer

Bachelor thesis
Supervisor: Christian Sillaber

University of Innsbruck

Contents

1 Introduction 1

2 Security Issues of the AUR 1

3 The Solution 3

3.1 Core Solution 3

3.2 Detailed Description 5

3.2.1 Blockchain 5

3.2.2 aursec-init 5

3.2.3 aursec — Architecture 6

3.2.4 aursec-hash 6

3.2.5 aursec-verify-hashes 7

3.2.6 aursec-chain 7

3.3 Terminal User Interface 7

3.4 Project Management 8

4 Things we Learned 8

5 Evaluation 9

List of Figures

1 Main Workflow 4

2 Decision Workflow 4

3 Timeline 8

List of Tables

1 Introduction

The Linux distribution Arch makes it easy to create custom packages and has a very active community. Because of this, there was the need for a place where users could upload their packages for others to use.

`ftp://ftp.archlinux.org/income` was created, where packages could be put until maintainers that were willing to do so could adopt them, but this delay was too long, so another solution was needed. The next improvement were the Trusted User Repositories, where some privileged users, which were many more than the maintainers of before, were allowed to host their own repositories for anyone to use.

The *Arch User Repository* were the natural evolution: By removing all middlemen, everyone can now upload their packages to one central place. [1]

The AUR is similar to PyPI (Python), NPM (Javascript) and Ruby gems, where all users can share their packages. They all share the problem that submitted packages are not necessarily audited or even checked by anyone.

2 Security Issues of the AUR

Ease of use appears to have been, if not the only, at least the primary design consideration of the Arch User Repository. This creates so many security issues that it is actually quite a task to think through all of them.

Local Package Creation

One of the most obvious problems is the installation procedure itself. The AUR doesn't host binary packages, which is a good thing. Instead, Arch packages are created locally from a bash file, the so-called `PKGBUILD`, containing metadata like name and version, the URLs and checksums of upstream sources, and functions for the compilation and packaging steps.

The AUR contains these `PKGBUILDS` and possible patches to be applied to the upstream sources in a git repository per package. A package file can be produced by cloning it's repository and using a tool called `makepkg` [2], which sources the script, downloads and verifies the sources, and calls the compilation and packaging functions.

This means that users can verify what they are compiling as opposed to blindly trusting binaries created by third-parties, but also that maintainers of AUR packages have a means of executing arbitrary shell commands on users' machines.

This is aggravated by the fact that `PKGBUILDS` can include a `.install` file into the built package, which will be executed *as root* when the package is actually installed. The risk also increases if so-called "AUR helpers" are used. These tools assist the users in installing packages from the AUR by automating the steps and behave like package managers. Some of them (notably `aurutils` [5], which is recommended by the authors) allow the users to inspect these files before continuing, but others are very unsafe in that they execute code before giving users the opportunity to inspect it, or incentivize them from doing so.

The Trust Issue

Another problem is that users are not given any reason to trust the maintainers. Unlike the official repositories, where maintainers are vetted, packages are (often manually) audited before being accepted, and everything must be signed with a trusted GNU Privacy Guard key, anyone can create an account and submit a new package to the AUR in a few minutes.

There is no admission procedure or audit system and no GnuPG web of trust in order to minimize the time needed to publish a package or update.

`makepkg` can verify GnuPG signatures for upstream sources, but the PKGBUILD itself could only be signed by using signed git commits, which is sadly not enforced or even officially recommended — and not supported by any AUR helper anyway.

Except when using the AUR helper `bauerbill` [7], which provides a basic user-side trust management system, the only way to be maintain reasonable trust is therefore to manually read every single file, which is cumbersome. Because only highly security-conscious users are willing to put in so much effort before trusting a PKGBUILD, most users are left vulnerable by the aforementioned issues.

Adopting orphan packages

The trust issue is made even worse by the fact that packages can silently and quickly be taken over by other maintainers due to the orphan/adoption system.

When a maintainer wants to stop maintaining a package, but the package is still useful and actively developed upstream, he has the option to *orphan* it rather than deleting it. Orphan packages can be *adopted* by any AUR user, at any time, without delay. This feature was designed to minimize update delay, which it does effectively; However, it also makes it easy for malicious agents to take over popular orphaned packages, manipulate them, and immediately orphan them afterwards.

The above Security issues lead to two concrete attack scenarios:

VCS Packages: Malicious Upstream

In some cases, the user is not adequately protected against malicious (or compromised) upstreams:

The AUR supports so-called *VCS packages* [3], which download sources from a version control system, such as Git or Mercurial, instead of downloading a fixed archive. This relieves the maintainer from updating his PKGBUILD for every new commit. `makepkg` will even automatically calculate the up-to-date version number using e.g. tags and commit numbers.

VCS packages were primarily designed to simplify the installation of up-to-date packages from source, and they do that very well; But they also introduce a big security issue: Since the PKGBUILD must not be updated between versions, it cannot contain checksums for the new version, either. This means that users don't have a way to verify the authenticity of the source that they are downloading, unless they can trust the upstream himself, meaning that no-one will notice if the upstream is compromised or makes malicious changes.

Tampered Packages: Malicious Maintainer

Users are also not adequately protected against malicious maintainers:

Because it's so easy to gain access to a package, e.g. by adopting an orphan or simply by creating it, and nothing is verified or audited before publication, It's easy for a malicious agent to modify a package. And because the PKGBUILD is not signed or hashed, users will not notice if the build instructions for a specific package version were modified. This allows targeted attacks:

If the time at which a target will update his machine is known and one has access to an AUR package which he is expected to update, malicious code can somehow be introduced into the corresponding PKGBUILD within that update window. This could be as simple as

changing the checksum if one also has access to the upstream source code (even a very careful user has no chance of noticing this attack), or executing innocuous code in the install file or PKGBUILD itself.

If the time frame is short, no other AUR user (and thus, *no-one*) would ever notice. One could only defend against this with a good local trust model, such as that possible with `bauerbill`, or manual cryptographic verification of the git commit to the AUR — assuming that the maintainer signs his commits, which is only very rarely the case.

3 The Solution

3.1 Core Solution

The solution for a few security-problems is a secure database, which contains hashes of versioned packages. Before a package is installed, the locale generated hash can be compared with the one in the database. This guarantee that the loaded package is the same as the package loaded by most of the other users.

To make the database as safe as possible, a blockchain is used. On this blockchain is a smart contract, which allows to call functions. With one of these functions it's possible to commit hashes of versioned packages. This hash will be saved in the blockchain if this user has not committed a other hash for the same versioned package before. Another function is used to get the current consensus hash and its number of commits of a versioned package. This is the first application which uses a blockchain to secure downloads.

Workflow

The following workflow is visualized by Figure 1.

1. First of all a `PKGBUILD` is downloaded and partially executed in a virtual area. Then this data get hashed.
2. The resulting local hash becomes compared with the current consensus hash of the versioned package of the blockchain [Figure 2].
3. Now the workflow splits into 3 different paths.
 - a) The hashes match and the number of commits is over the threshold or the user decides to trust the local generated test. The package is created and installed. (*Followed by step 4.*)
 - b) The hashes don't match and/or the number is below the threshold but the user want to create and install the package. (*End of the workflow.*)
 - c) The hashes don't match and/or the number is below the threshold and the user doesn't want to create the package. (*End of the workflow.*)
4. The local hash is committed to the blockchain (this is a transaction).
5. All nodes of the blockchain-network get the transaction.
6. The transaction is contained in the next mined block.
7. The block is added to the blockchain.




Figure 1: Main Workflow




Figure 2: Decision Workflow

3.2 Detailed Description

3.2.1 Blockchain

general

Workflow:

Smart Contracts:

specific

Our blockchain has to fulfill several requirements:

- Since we target Arch Linux, it should be easy to install on this platform.
- It needs to provide an API which allows external scripts to work with it.
- It needs to provide smart contracts to maintain the constraints of the hash storage.
- It needs to provide an easy way to create private networks separate from the main one, because full networks are huge.

The blockchain of choice was Ethereum because it is the only production-ready infrastructure providing smart contracts. It even allows them to be written in a high-level language, *Solidity*, which makes them straightforward to write and understand.

smart contracts Aursec has two smart contracts. The first one is a formal contract, which allows the owner of the contract to delete the contract. The second contract is a child of the first. In this contract are the “user-functions”, which allows all users to send hash-commits and request the current consensus hash of a versioned package. The Contract allows a user to commit one hash per versioned package. Further commits of hashes of the same versioned package will not be considered. The current consensus of a versioned package is the most often committed hash so the hash of a package is not fixed by the first commit.

network Like other blockchains, the Ethereum network is peer-2-peer, but private networks need to provide their own *bootnode* through which peers announce themselves. Our bootnode is active 24/7 on a DigitalOcean droplet provided by our supervisor.

interfaces The Ethereum blockchain provides 2 different interfaces, the IPC (interprocess communication), which provides an interactive javascript shell, and the HTTP RPC (remote procedure call) interface. The IPC interface of our blockchain is deactivated because users don’t need to manually use the javascript shell: the only required interaction with the blockchain is through the shell script `aursec-chain` 3.2.6 and the Python script `aursec-tui` 3.3. These scripts communicate with the blockchain through the RPC interface.

3.2.2 aursec-init

Aursec-init is a shellscript which allows the user to create all requirements just by running it. It also allows to overwrite an existing aursec-blockchain with a new one.

workflow of the initialization:

1. Create needed folders and markers (markers are needed by TODO)
2. Create the blockchain from our genesis block.
3. Create a new user with a random password which is saved in a file.
4. Create the *Directed Acyclic Graph* for the blockchain, which is a 1GB dataset. The DAG is needed for mining new blocks. [6]
5. Set safe permissions for the above files and folders.
6. Mine a few blocks to trigger the synchronization and to have enough ether to be able to commit a few hashes from the start.

3.2.3 aursec — Architecture

We designed `aursec` to be a good UNIX citizen:

- Modular design
multiple small tools doing one thing and doing it well
- Adherence to the universal interface [4]
Work on streams of text on `stdin` and `stdout`
- Written in Bash and using existing tools where possible
- Maximize concurrency using a pipeline

`aursec` builds a pipeline of two other tools, `aursec-hash` and `aursec-verify-hashes`, which produces lines containing the id and hash of a package as well as the hash representing the current consensus on the blockchain and the number of times that hash was submitted. It then inserts itself into that pipeline and iterates over the lines of items using a *while-read* loop and traverses the aforementioned state machine for each item.

This architecture has several advantages: It is straightforward to understand because it follows standard UNIX patterns, which also makes it very maintainable. The free 3-level parallelism gained by the pipeline is a very useful advantage in itself, but even more so because all 3 tools are primarily I/O-bound: `aursec-hash` reads and hashes lots of files, `aursec-verify-hashes` constantly queries (and waits for) the blockchain, and `aursec` tends to spend much time waiting for user input. Thus, the concurrency is even more important because it allows work to continue in the background while `aursec` waits for the user. Indeed, the background tasks tend to be finished in most practical situations before the user has had time to inspect the second or third warning.

3.2.4 aursec-hash

`aursec-hash` has the simple task of producing an ID (`$pkgname-$pkgver-$pkgrel`) and a hash from PKGBUILDs.

The id could be parsed from the `.SRCINFO`, which is a plain text file. But VCS packages don't have up-to-date version information in their PKGBUILD, which means that it must be interpreted; This is annoying, but we only source the PKGBUILD in a `firejail` sandbox to minimize the inherent risk of executing foreign turing-complete code. This allows us to

get an accurate ID for VCS packages, but also to include the actual sources in the hash, thereby compensating for the lack of hashes in the PKGBUILD of VCS packages.

The PKGBUILD is hashed after stripping its comments, and VCS sources, if they exist, are hashed using `find`. Finally, all hashes are combined by another call to the hash command. Currently, `sha256sum` is used for its good speed and security.

3.2.5 aursec-verify-hashes

This tool is not very interesting; It simply fetches the current consensus for every package ID on `stdin`, computes whether it matches with the locally computed hash, and appends that data to the output stream. Doing this in a separate pipeline step is very worth it because requests from the blockchain are comparatively slow.

3.2.6 aursec-chain

aursec-chain is a shellscript which allows the user to intergate with the blockchain. The script itself communicates with the blockchain through JSON RPC. To provide the user all needed commands, the script has different arguments.

mine This argument needs more arguments to work.

- **start**: Starts mining.
- **stop**: Stops mining.
- **N blocks**: Wait until mining is stopped and then mines N blocks.
- **auto**: This command is only used by a systemd timer to mine blocks periodically.

commit-hash This argument needs two more arguments to work. The first one has to be the versioned package id and the second has to be the locally generated hash of the package. The script then sends a transaction to the blockchain (if enough ether is available). This transaction will be verified by the next mined block.

get-hash This argument needs one more argument to work. This argument has to be the versioned package id. Then *aursec-chain* calls a method which returns the current consensus hash and the number of commits of this hash.

3.3 Terminal User Interface

aursec-tui is a Python script. The TUI gives a overview over all mined blocks, their hashes, the miner of the block and eventual transactions which are saved in the block.

It is split in two parts. The first part is needed to gain the data from the blockchain and save it, the other part formats the information and displays it. To display the results as fast as possible a thread searches the blockchain in the background. Any additional data will be displayed after the next refresh. The script offers the user two settings for filtering the results.

- **only mine**: only display blocks which where mined by the user himself.
- **only transactions**: only display blocks with transactions.

The settings can be combined with the result that all blocks mined by the user with transactions will be displayed.

3.4 Project Management

This project has involved three people. Two students who implemented the tool and their superior. So there had to be a lot of communication. The communication tools were offline meetings, Slack and emails. For files sharing Github is used.

timeline

From the timeline (Figure 3) can be derived that this project was realized in less then 9 months. The real worktime exposure ratio of the three main-task (reading, programming and writing) differs extremely from the visualized one. In reality we read 50%, programmed 40% and wrote 10% of the worktime. It is also visible, that we didn't manage to complete three milestones in time. This was caused by different circumstances, but we managed to complete the hole project in time cause we adjusted our planing every few weeks.




Figure 3: Timeline

4 Things we Learned

Ethereum-Solidity

JSON RPC

The JSON RPC is used by the aursec-ethereum-blockchain for remote access to the block. Methods can be called or transactions can be sent by sending a JSON-Object to the block chain. The answer from the block chain is also always a JSON-Object. One big advantage is the easy parsing from and to JSON-Objects both in Bash and Python. The best example is visible in the code of aursec-tui (3.3).

Bash

Python-Urwid

Urwid is a *Terminal User Interface* (TUI) library for Python. It provides ready-made widgets which make it easy to create structured user interfaces.

5 Evaluation

References

- [1] ArchWiki. Arch User Repository — Arch Wiki. https://wiki.archlinux.de/title/Arch_User_Repository, 2017. accessed March 17, 2017.
- [2] ArchWiki. Creating packages — Arch Wiki. https://wiki.archlinux.org/index.php/Creating_packages, 2017. accessed March 18, 2017.
- [3] ArchWiki. VCS Package Guidelines — Arch Wiki. https://wiki.archlinux.org/index.php/VCS_package_guidelines, 2017. accessed March 18, 2017.
- [4] Peter H. Salus. *A Quarter Century of UNIX*. ACM Press/Addison-Wesley Publishing Co., New York, NY, USA, 1994.
- [5] Alad Wenter. aurutils: helper tools for the AUR. <https://github.com/AladW/aurutils>, 2016-2017. accessed March 18, 2017.
- [6] Ethereum Wiki. Ethash-DAG — Ethereum Wiki. <https://github.com/ethereum/wiki/wiki/Ethash-DAG>, 2017. accessed March 20, 2017.
- [7] xyne. Bauerbill: Extension of Powerpill with AUR and ABS support. <http://www.xyne.archlinux.ca/projects/bauerbill/>, 2015-2017. accessed March 18, 2017.